

Computer Arithmetic

**What is the difference between
int and unsigned int?**

Pat Hanrahan

cs107e

Addition

Binary Addition

	Carry
00000111 A	
+00001011 B	

	Sum

Binary Addition

	1	Carry
00000111	A	
+00001011	B	

	0	Sum

Binary Addition

	11	Carry
00000111	A	
+00001011	B	

	10	Sum

Binary Addition

00001111 Carry

00000111 A

+00001011 B

00010010 Sum

Binary Addition

```
11111111  Carry
 11111111  A
+00000001  B
-----
100000000  Sum
```

To represent the result of adding two n -bit numbers with full precision requires $n+1$ bits


But we only have 8-bits!

```
sum = 0b00000000 = (A+B)%256
```

9 9 9 9 9 9

9 9 9 9 9 9

0 0 0 0 0 0



Add 2 1-bit numbers

a	b	sum
---	---	-----

0	0	00
---	---	----

0	1	01
---	---	----

1	0	01
---	---	----


1	1	10
---	---	----

Add 2 1-bit numbers (Half Adder)

a	b	sum
0	0	00
0	1	01
1	0	01
1	1	10

bit 0 of sum: $S = a \oplus b$

bit 1 of sum: $C = a \& b$




Have reduced addition to logical operations!

Add 3 1-bit numbers

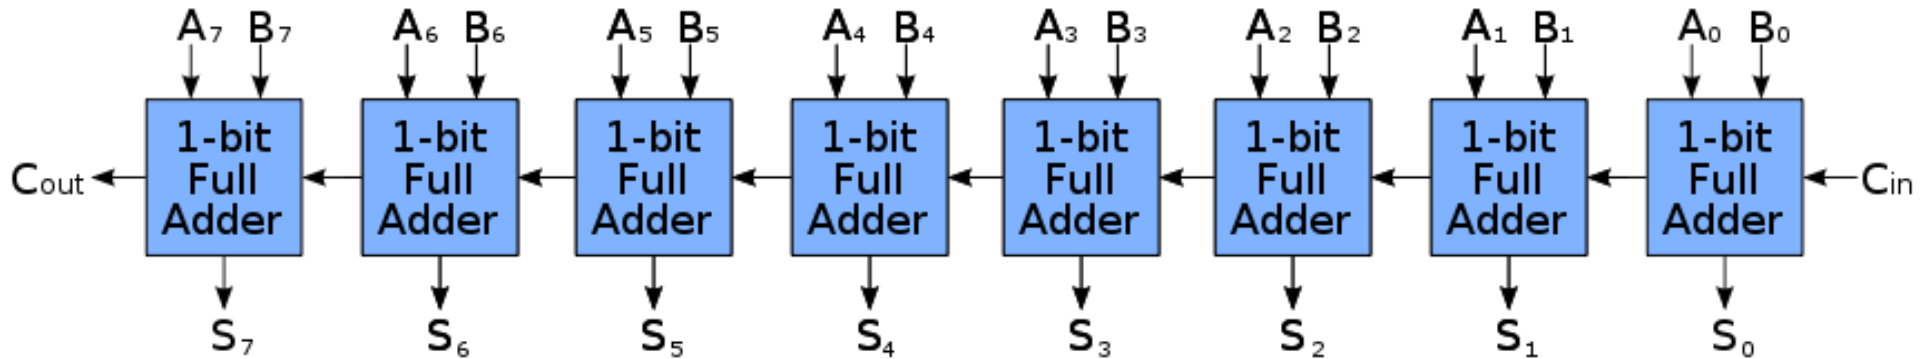
a	b	c	=	c	s
0	0	0		0	0
0	1	0		0	1
1	0	0		0	1
1	1	0		1	0
0	0	1		0	1
0	1	1		1	0
1	0	1		1	0
1	1	1		1	1

Add 3 1-bit numbers (Full Adder)

a	b	ci	=	co	s
0	0	0		0	0
0	1	0		0	1
1	0	0		0	1
1	1	0		1	0
0	0	1		0	1
0	1	1		1	0
1	0	1		1	0
1	1	1		1	1



8-bit Ripple Adder



Note C_{in} and C_{out}

```
// Multiple precision addition  
// http://godbolt.org/g/HMYrme
```

```
uint64_t add64(uint64_t a, uint64_t b)  
{  
    return a + b;  
}
```

```
add64:  
    adds r0, r0, r2  
    adc  r1, r1, r3  
    bx   lr
```

Subtraction

Binary Subtraction

00000001 Borrow

00000110 A

-00000001 B

00000101 Sub

Do we need to build subtraction hardware?

BIG IDEA: Define subtraction using addition


A clever way of defining subtraction by 1 is to find a number to add that yields the same result as the subtract by 1.

This number is the *negative* of the number.

More precisely, this number is the number that when added to 1, results in 0 (mod 16)

$$0x1 - 0x1 = 0x1 + 0xf = 0x10 \% 16 = 0x0$$

0xf can be *interpreted* as -1



Signed 4-bit numbers,

0x0 = 0

0xf = -1

0xe = -2

...

0x8 = -8 (could be interpreted as 8)

0x7 = 7

...

0x1 = 1

0x0 = 0

if we choose to *interpret* 0x8 as -8,
then the most-significant bit of the
number indicates that it is negative (n)

signed int vs as unsigned int

**Are just different interpretations of
the bits comprising the number**

0xff vs -1

Negation

How do we negate an 8-bit number?

Subtract it from 2^8 (0b100000000)

$-x = 0b100000000 - x$ (two's complement)

Since then $(x + (-x)) \% 256 = 0$

11111111	Borrow
100000000	
-00000001	

11111111	

100000000	Carry
00000001	
+11111111	

00000000	

Another way

Rewrite $100000000 = (11111111 + 1)$

$$\begin{aligned} -x &= (11111111+1)-x \\ &= (11111111-x)+1 \quad (\text{one's complement}) \\ &= \sim x + 1 \end{aligned}$$

E.g. $-1 = 0b11111111$

$\sim 00000001 = 11111110$ (\sim is invert)

$$11111110 + 00000001 = 11111111$$

Subtraction is converted negation + addition

-B is implemented using $\sim B + 1$

$$A - B = A + \sim B + 1$$

$$01 - 00 = 01 + ff + 01 = 01 + c$$

$$01 - 01 = 01 + fe + 01 = 00 + c$$


$$01 - 02 = 01 + fd + 01 = ff$$

Note the carry out bit c

The +1 can be done by setting Cin to 1

**Addition, Subtraction, and Negation
of signed and unsigned numbers
are the same!**

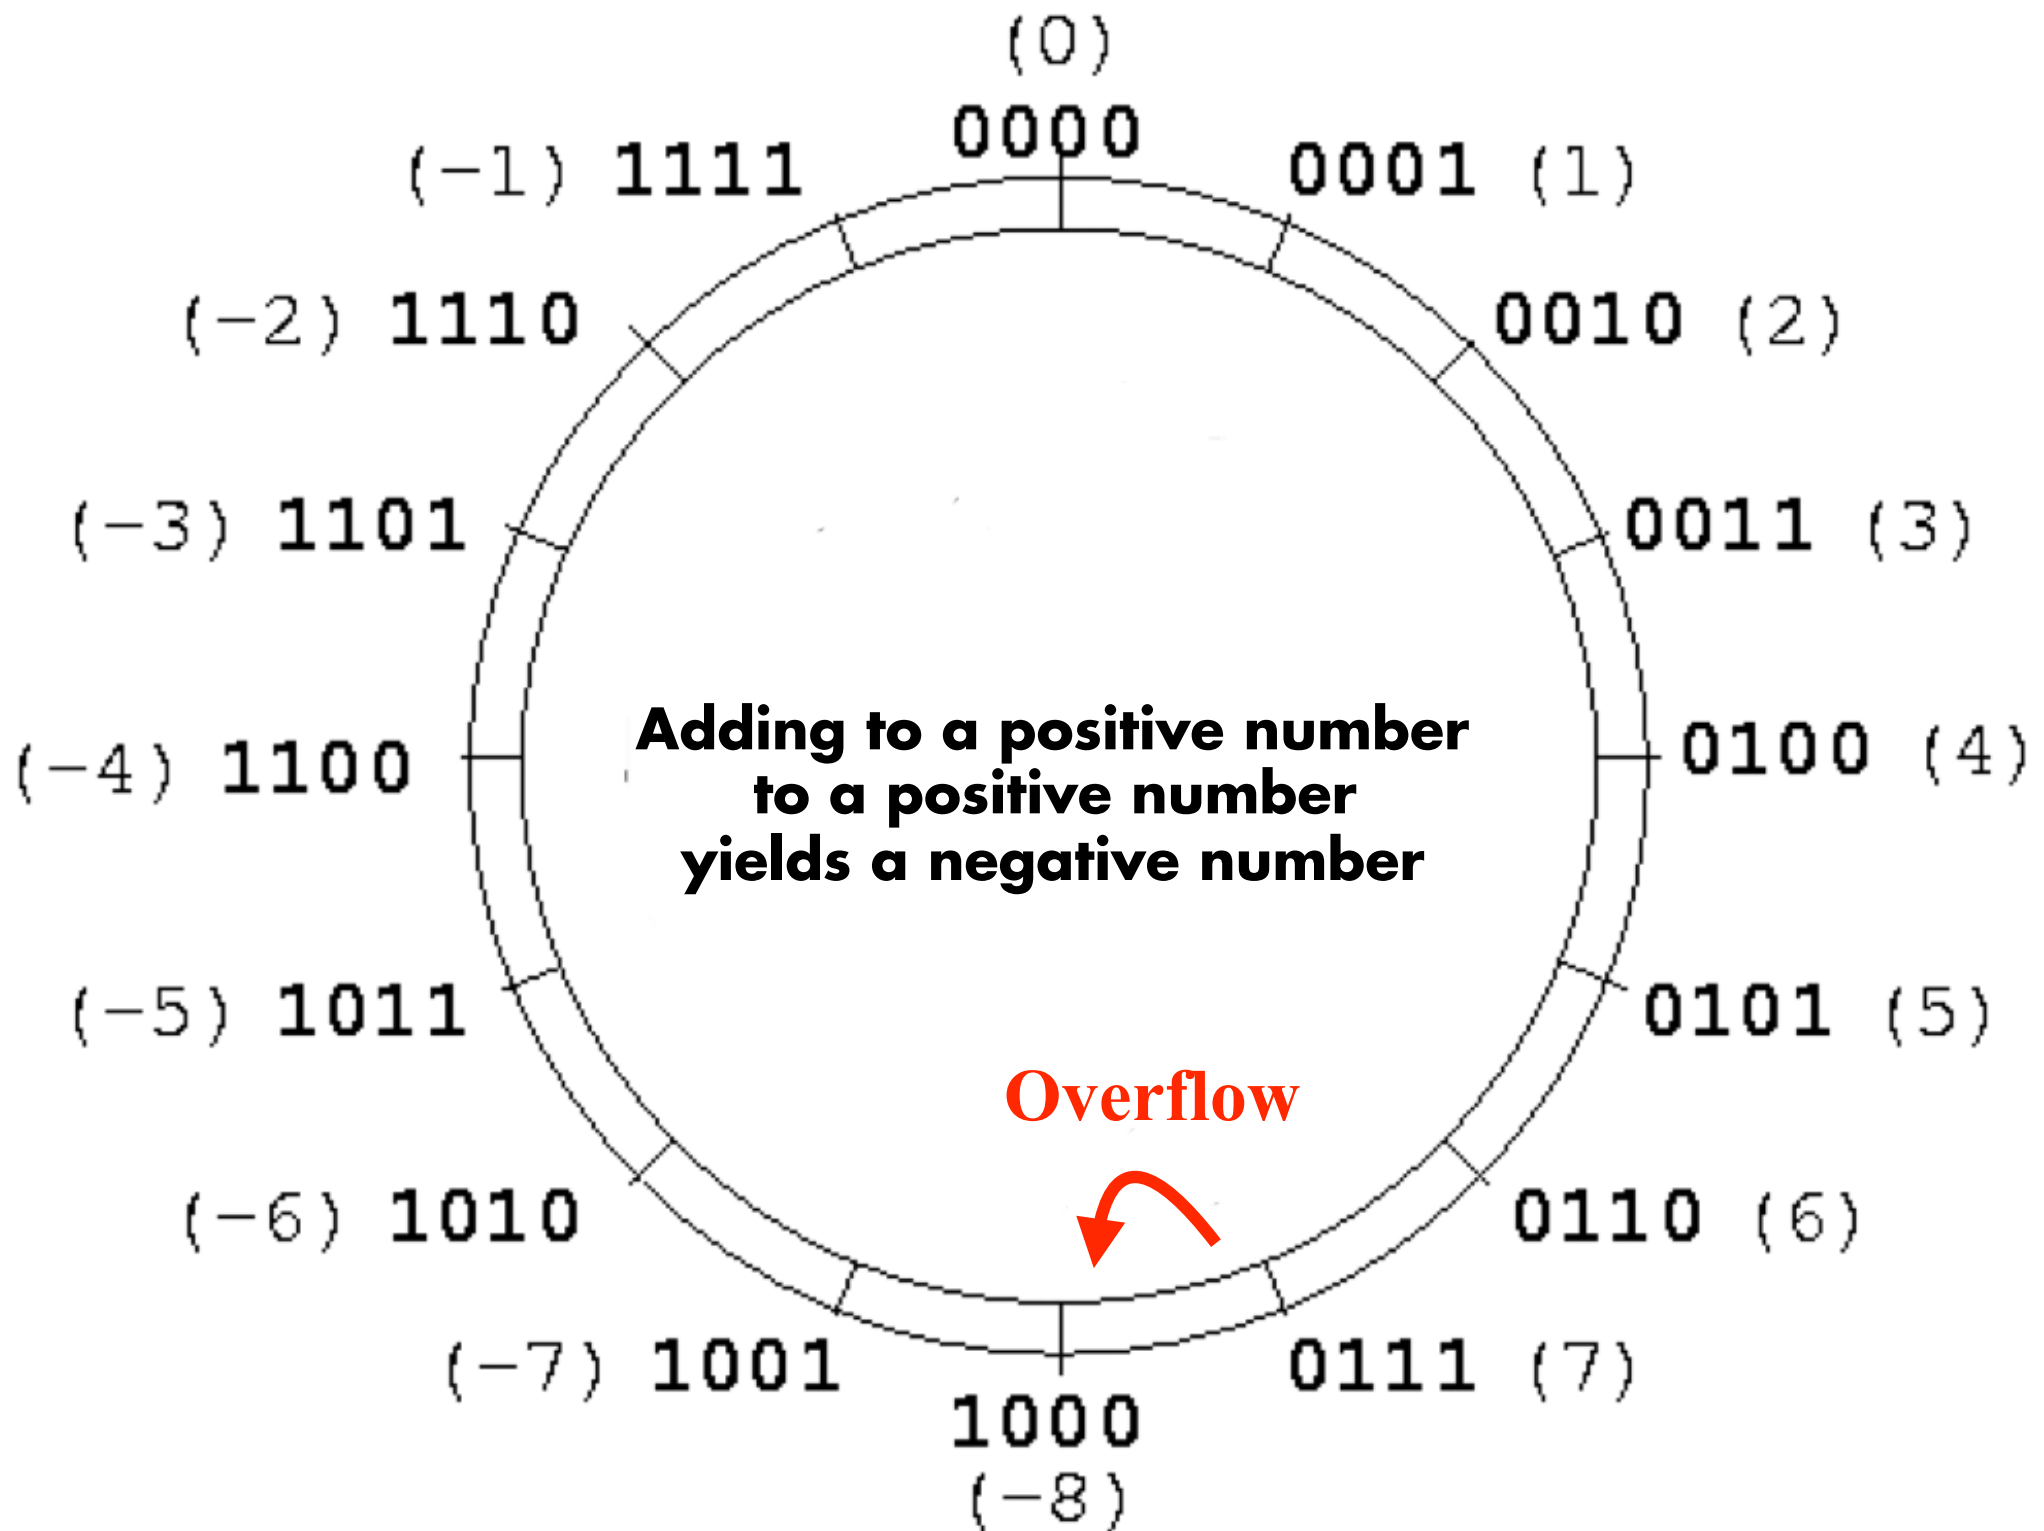
CPSR




Arithmetic instructions set N, Z, C, V

Logic instructions just set N, Z


What is V?

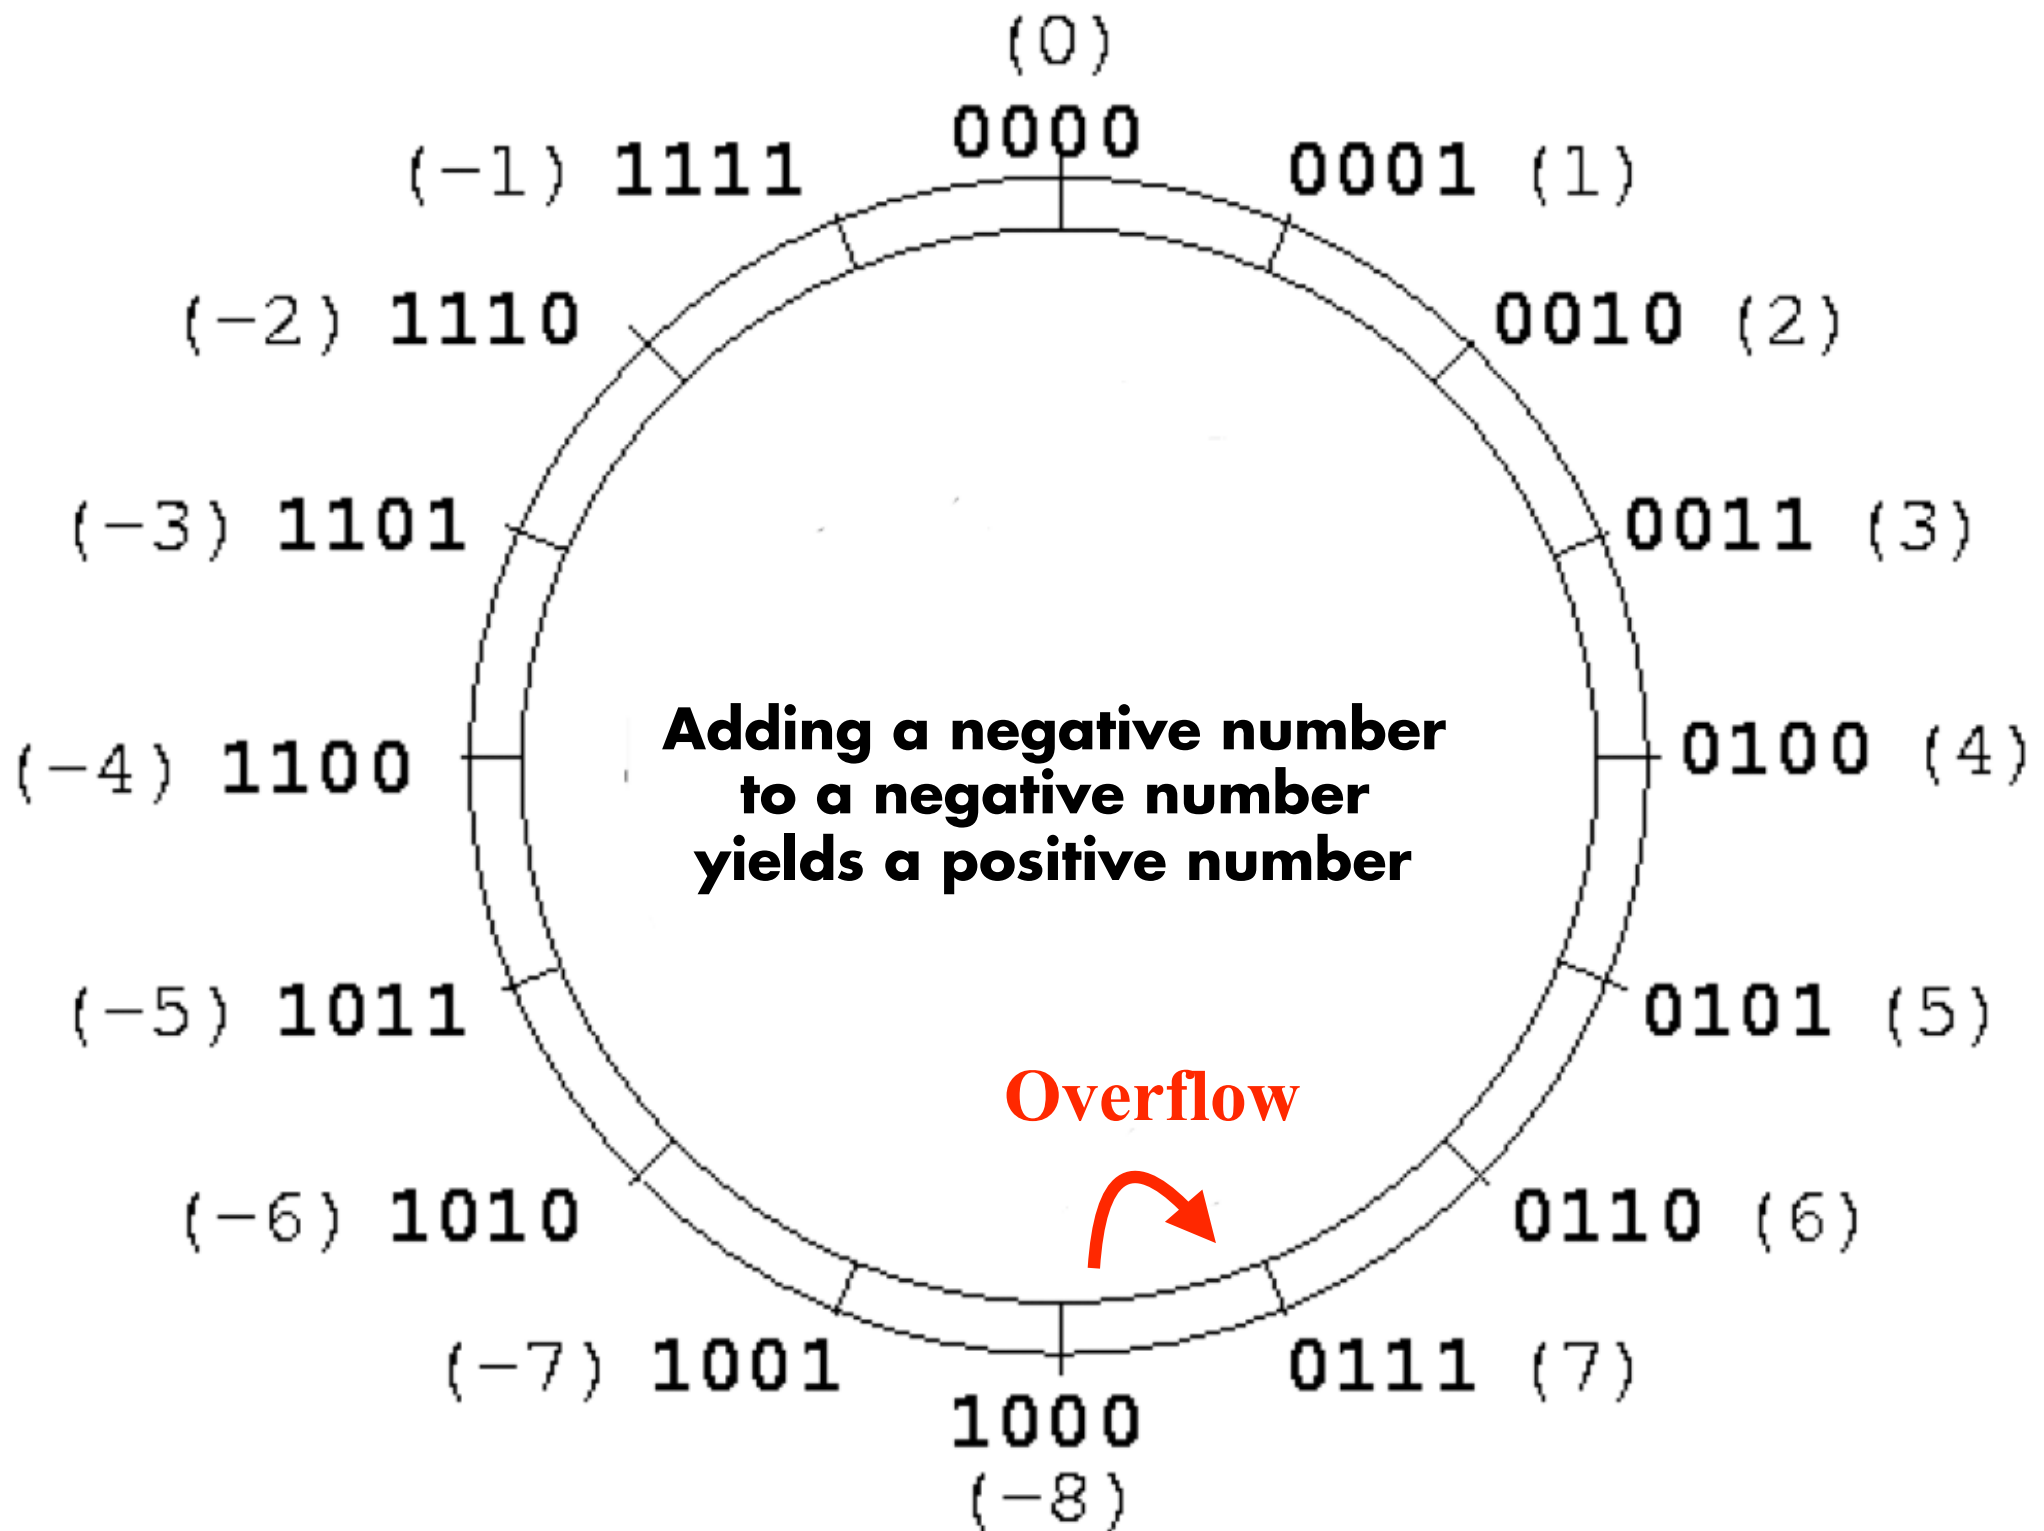



- In two's complement, when you exceed the maximum value of int (2,147,483,647), you “wrap around” to negative numbers:



- Here is the link after Google upgraded to 64-bit integers:







Comparison

=

Subtract and Look at Flags

Code	Suffix	Flags	Meaning
0000	EQ	Z set	equal
0001	NE	Z clear	not equal
0010	CS	C set	unsigned higher or same
0011	CC	C clear	unsigned lower
0100	MI	N set	negative
0101	PL	N clear	positive or zero
0110	VS	V set	overflow
0111	VC	V clear	no overflow
1000	HI	C set and Z clear	unsigned higher
1001	LS	C clear or Z set	unsigned lower or same
1010	GE	N equals V	greater or equal
1011	LT	N not equal to V	less than
1100	GT	Z clear AND (N equals V)	greater than
1101	LE	Z set OR (N not equal to V)	less than or equal
1110	AL	(ignored)	always

**Methods used to compare
signed and unsigned numbers
are NOT the same!**

Type Conversion

Jedi Job

Interview Questions

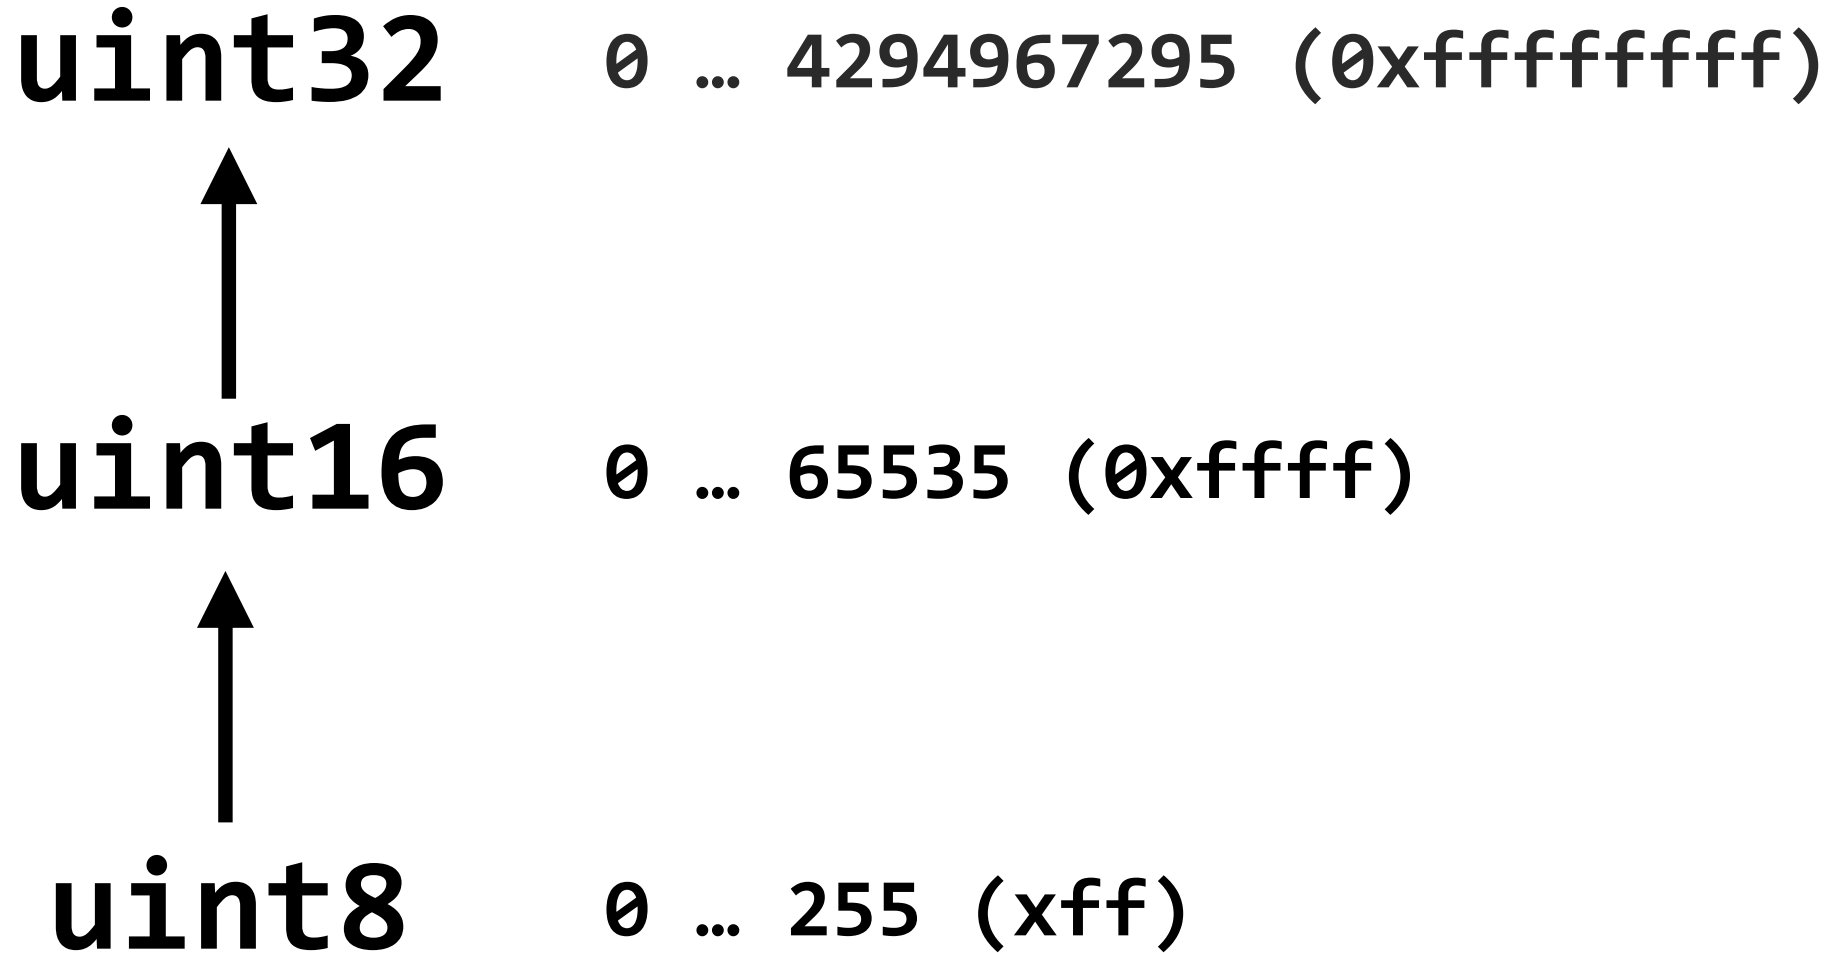
```
#include <stdint.h>
```

```
uint16_t x = 0xffff;
```

```
uint32_t y = x;
```

```
// y?
```

Type Hierarchy



Types are *sets* of allowed values

Arrow indicate *subsets*: $\text{uint16} \subset \text{uint32}$

uint32



uint16



uint8

**Type Conversion is Safe
(values preserved)**

```
#include <stdint.h>
```

```
uint16_t x = 0xffff;
```

```
uint32_t y = x;
```

```
// y = 0x0000ffff
```

```
int16_t x = -1;
```


```
int32_t y = x;
```

```
// y?
```


int32 -2,147,483,648 ... 2,147,483,647



int16 -32768 ... 32767



int8 -128 ... 127

Type Conversion is Safe
(values preserved)

What is the maximum value for a int32?



606

integer



share improve this question



92

I can never remember that number. I need a memory rule.

edited May 28 '14 at 14:09



Ben Hoffstein

49.5k ● 5 ● 66 ● 101

asked Sep 18 '08 at 17:18



Flinkman

5,181 ● 4 ● 18 ● 48

107 Why would you need the exact number? I remember " $(2^{31})-1$ " or " ± 2 billion" and that's good enough for everything I ever needed. — Joachim Sauer Mar 3 '09 at 11:21

27 unsigned: $2^{32}-1 = 4 \cdot 1024^3 - 1$; signed: $-2^{31} \dots +2^{31}-1$, because the sign-bit is the highest bit. Just learn $2^0=1$ to $2^{10}=1024$ and combine. $1024=1k$, $1024^2=1M$, $1024^3=1G$ — comonad Mar 28 '11 at 20:01

6 I generally remember that every 3 bits is about a decimal digit. This gets me to the right order of magnitude: 32 bits is 10 digits. — Barnar Oct 2 '13 at 15:11

30 Answers

active

oldest

votes



2397

share improve this answer



It's 2,147,483,647. Easiest way to memorize it is via a tattoo.

edited Oct 20 '14 at 16:30



Allbite

1,415 ● 1 ● 13 ● 15

answered Sep 18 '08 at 17:20



Ben Hoffstein

49.5k ● 5 ● 66 ● 101

```
int16_t x = -1;  
int32_t y = x;
```

```
// what is the value of y?
```

```
// x = -1 = 0xffff  
// y = -1 = 0xffffffff
```

```
// Sign extension
```

```
int8_t 0xfe -> int32_t 0xfffffffffe
```

```
int8_t 0x7e -> int32_t 0x0000007e
```

// Assembly language

LSL r0,r0,#24

ASR r0,r0,#24

fe000000 1 1 1 1 1 1 1 0

asr

The diagram shows a horizontal row of 32 square boxes, each containing a single digit. The first 8 boxes contain the digit '1', and the remaining 24 boxes contain the digit '0'. To the left of the first box is the label 'asr'. Above the first box is a large circular arrow pointing clockwise, indicating a right shift operation.

[illegible]

// Sign extend instructions: sxth and sxth

```
int32_t x = -1;
```

```
int16_t y = x;
```

```
// y?
```

```
int32_t x = -1;
```

```
int16_t y = x;
```

```
// y = -1
```

```
int32_t x = 4*INT16_MIN;
```

```
int16_t y = x;
```

```
// y?
```

```
int32_t x = 4*INT16_MIN;  
int16_t y = x;
```

```
// y = 0
```

 value has changed

int32



int16



int8

Defined (remove most significant bits)

Dangerous (doesn't preserve all values)

```
int32_t  x = -1;
```

```
uint32_t y =  x;
```

```
// y?
```

```
int32_t  x = -1;  
uint32_t y =  x;
```

```
// y = 0xffffffff
```

 value has changed

(y is a large positive number!)

uint32 ← int32

uint16 ← int16

uint8 ← int8

Defined (copies bits)

uint32 ← int32

uint16 ← int16

uint8 ← int8

Dangerous! (neg become large)

uint32 → int32

uint16 → int16

uint8 → int8

**Technically Not Defined
(arm: copies bits)**

uint32



int32

uint16


int16

uint8

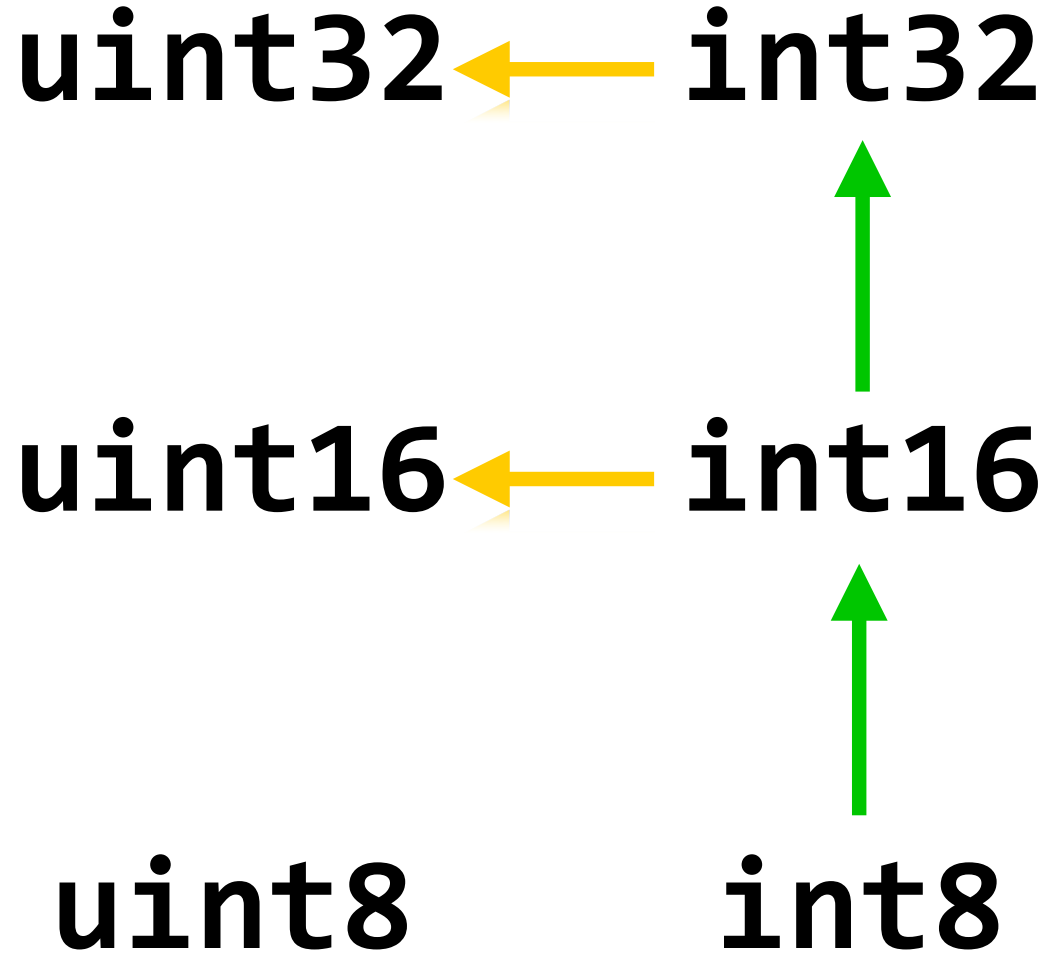
int8



Safe?



Safe?



Defined, Dangerous

Binary Operations

Type promotions for binary operations

Note that the type of the result can be different than the type of the operands!

	u8	u16	u32	u64	i8	i16	i32	i64
u8	i32	i32	u32	u64	i32	i32	i32	i64
u16	i32	i32	u32	u64	i32	i32	i32	i64
u32	u32	u32	u32	u64	u32	u32	u32	i64
u64	u64	u64	u64	u64	u64	u64	u64	u64
i8	i32	i32	u32	u64	i32	i32	i32	i64
i16	i32	i32	u32	u64	i32	i32	i32	i64
i32	i32	i32	u32	u64	i32	i32	i32	i64
i64	i64	i64	i64	u64	i64	i64	i64	i64

arm-none-eabi-gcc type promotions

A collection of black ant icons scattered across the page. There are 10 ants in total, some at the top, some in the middle, and some at the bottom left. Each ant is a simple black silhouette with six legs and two antennae. The text "Bugs, Bugs, Bugs" is centered in the middle of the page in a bold, black, sans-serif font.

Bugs, Bugs, Bugs

```
#include <stdio.h>
```

```
int main(void)
```

```
{
```

```
    int a = -20;
```

```
    unsigned int b = 6;
```

```
    if( a < b )
```

```
        printf("-20<6 - all is well\n");
```

```
    else
```

```
        printf("-20>=6 - omg \n");
```

```
}
```

**Whenever you mix
signed and unsigned numbers
you get in trouble**

Bjarne Stroustrup

Summary

Negation is performed by forming the two's complement

- **Signed numbers are represented in two's complement ($-x = 2^n - x = \sim x + 1$)**

In 2's complement,

- **Arithmetic between signed and unsigned numbers is identical**
- **Comparison between signed and unsigned numbers is different**

Know the rules for type conversion, watch out for implicit type conversions

C Type Conversion and Promotion Rules

The semantics of numeric casts are:

Casting from a larger integer to a smaller integer (e.g. u32 -> u8) will truncate

Casting from a smaller integer to a larger integer (e.g. u8 -> u32) will zero-extend if the source is unsigned sign-extend if the source is signed

Casting between two integers of the same size (e.g. i32 -> u32) is a no-op

6.3.1.3 Signed and unsigned integers conversions

1 When a value with integer type is converted to another integer type, if the value can be represented by the new type, it is unchanged.

2 Otherwise, if the new type is unsigned, the value is converted by repeatedly adding or subtracting one more than the maximum value that can be represented in the new type until the value is in the range of the new type.

3 Otherwise, if the new type is signed and the value cannot be represented in it; either the result is implementation-defined or an implementation-defined signal is raised.

6.3.1.8 Usual arithmetic conversions

1 If both operands have the same type, then no further conversion is needed.

2 Otherwise, if both operands have signed integer types or both have unsigned integer types, the operand with the type of lesser integer conversion rank is converted to the type of the operand with greater rank.

3 Otherwise, if the operand that has unsigned integer type has rank greater or equal to the rank of the type of the other operand, then the operand with signed integer type is converted to the type of the operand with unsigned integer type.

4 Otherwise, if the type of the operand with signed integer type can represent all of the values of the type of the operand with unsigned integer type, then the operand with unsigned integer type is converted to the type of the operand with signed integer type.

5 Otherwise, both operands are converted to the unsigned integer type corresponding to the type of the operand with signed integer type.

```
unsigned int timer_get_ticks(void)
{
    return *SYSTIMERCLO;
}
```

```
void timer_delay_us(unsigned int usecs)
{
    unsigned int start=timer_get_ticks();
    while (timer_get_ticks()-start < usecs)
        { /* spin */ }
}
```

// The timer continuously ticks.

// Does this code work if the timer overflows?