

Goals for today



Build process

How do source files become an executable?

What does the linker do?

Diagnosing build errors

Modules and libraries

What makes for good design?

Admin

Keep eye on piazza!

Strategies for assign3

Where are we going?

Processor and memory architecture

Peripherals: GPIO, timers, UART

Assembly language and machine code

From C to assembly language

Function calls and stack frames

Serial communication and strings

Modules and libraries: Building and linking

Memory management: Memory map & heap

gpio
timer
uart
strings
printf
malloc
keyboard
fb
gl
console
shell



Good modules

Meaningful decomposition of the system
into parts (modules)

Interfaces and implementations

- Provide easy-to-use interface to users
- Hide obscure details of implementation

Tested independently with unit tests

Obi-wan, how should I design my modules?

"The Build"



```
APPLICATION = clock
MY_MODULES = timer.o gpio.o

CFLAGS = -Og -g -Wall -std=c99 -ffreestanding
LDFLAGS = -nostdlib -T memmap
LDLIBS =

all : $(APPLICATION).bin

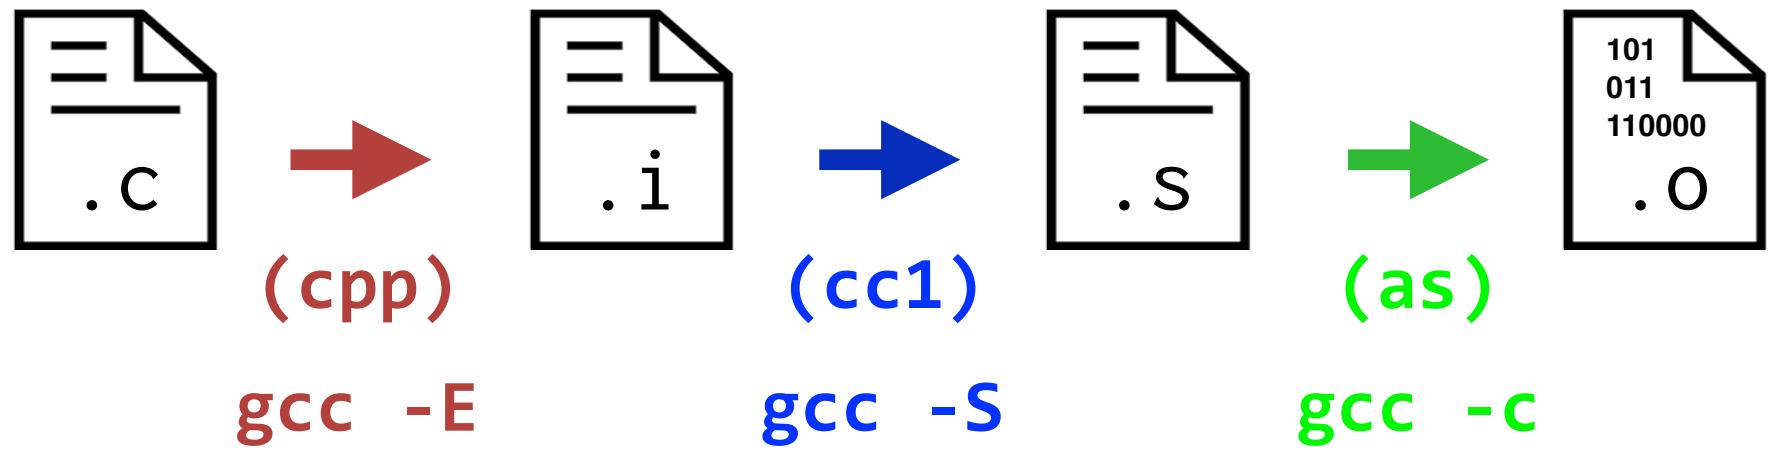
%.bin: %.elf
    arm-none-eabi-objcopy $< -O binary $@

%.elf: %.o $(MY_MODULES) start.o cstart.o
    arm-none-eabi-gcc $(LDFLAGS) $^ $(LDLIBS) -o $@

%.o: %.c
    arm-none-eabi-gcc $(CFLAGS) -c $< -o $@

%.o: %.s
    arm-none-eabi-as $< -o $@
```

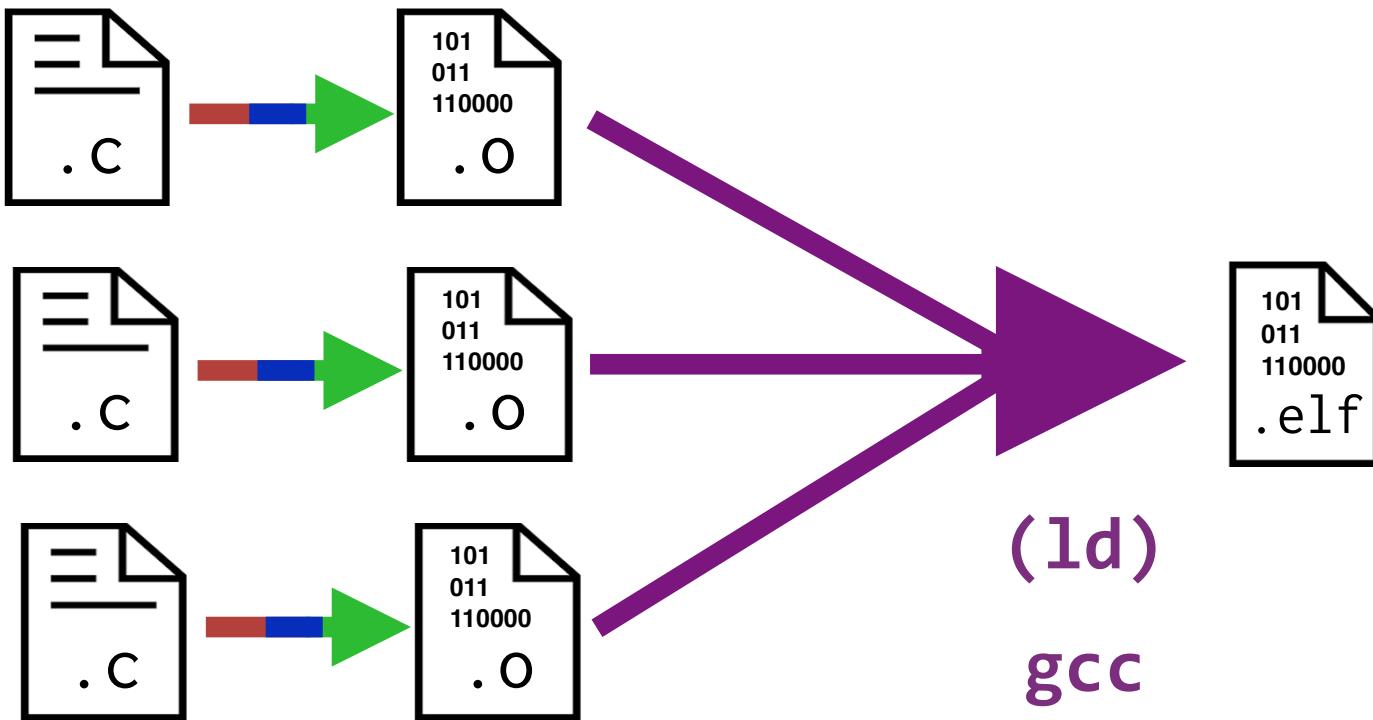
Compiling a single module



want to see intermediate files?

gcc --save-temp

Linking modules into executable



```
gcc clock.o timer.o gpio.o start.o cstart.o
```

What does the linker do?

Combines one or more object files into executable

Resolve inter-module references

Consolidate code, data sections across all modules

Arrange sections in output file in proper order

Symbols

Single global namespace

A name can have only one definition

Need conventions to avoid clashes

e.g. `gpio_init` versus `timer_init`

Qualifier dictates visibility, linking

static private to module, not visible outside, no linking needed

extern public, exported for use by other modules

Linker only concerned with `extern` symbols

Symbol resolution

Rule: Every symbol referenced must be defined
once and exactly once

Possible errors:

Multiply-defined: two definitions for same symbol

Undefined: needed symbol never defined

Linker resolution process



Set D (symbols that have been defined)

Set U (symbols that have been referenced, but not yet defined)

Linker processes .o files in order from left to right

For each .o, updates **Set U** and **Set D** and goes on to next
until all processed

Link successful: Set U empty and set D contains no duplicates!



Libraries

An archive .a is just a collection of modules (.o files)

The linker scans the library to find definitions for symbols in U (the undefined set). When a symbol is found, the entire module and all its functions are linked in.

If adding that module results in more undefined symbols; linker searches library for definition of these symbols and links those modules; repeats until no more definitions of undefined symbols are found, linker moves on to next library

Diagnosing build errors

What tool detects/reports... ?

mismatched type

function call with wrong arguments

forgot #include

forgot to link with library

use variable uninitialized

#include "wrongname.h"

two functions with same name

Building with libpi

```
APPLICATION = main
MY_MODULES = printf.o strings.o

CS107E = ../cs107e.github.io/cs107e
CFLAGS = -I$(CS107E)/include -g -Wall -Og -std=c99 -ffreestanding
LDFLAGS = -nostdlib -T memmap -L$(CS107E)/lib
LDLIBS = -lc_pi -lpi -lgcc

all : $(APPLICATION).bin $(MY_MODULES)

%.bin: %.elf
    arm-none-eabi-objcopy $< -O binary $@

%.elf: %.o $(MY_MODULES) start.o cstart.o
    arm-none-eabi-gcc $(LDFLAGS) $^ $(LDLIBS) -o $@

...
```

Combining multiple modules (.o)
into a single executable (.elf)

memmap

```
SECTIONS
{
    .text 0x8000 : { start.o(.text*)
                      *(.text*) }
}
```

Why must **start.o** go first?

What is the significance of **0x8000**?

```
% arm-none-eabi-nm -n clock.elf
00008000 T _start
0000800c t hang
00008010 T square
0000801c T blink
00008070 T main
0000808c T timer_init
00008090 T timer_get_ticks
00008098 T timer_delay_us
000080a4 T timer_delay_ms
000080c0 T timer_delay
000080e0 T gpio_init
000080e4 T gpio_set_function
000080e8 T gpio_get_function
000080f0 T gpio_set_input
000080f4 T gpio_set_output
000080f8 T gpio_write
000080fc T gpio_read
00008104 T _cstart
```

size reports the size of the segments:

% arm-none-eabi-size				filename
text	data	bss	dec	hex
340	0	0	340	154 clock.elf
% arm-none-eabi-size *.o				
text	data	bss	dec	hex filename
124	0	0	124	7c clock.o
80	0	0	80	50 cstart.o
36	0	0	36	24 gpio.o
16	0	0	16	10 start.o
84	0	0	84	54 timer.o

Note sum of sizes for .o's equals size of main.elf

Relocation

Relocation



In the .o:

each symbol assigned a temporary address (offset within module)
access to other symbols in same module is PC-relative

In the executable:

all modules consolidated
first module is laid down at fixed address (0x8000)
base for each module is offset by size of previous modules
final address of each symbol =

module_base + symbol_offset_within_module

```
// start.s
```

```
.globl _start
_start:
    mov sp, #0x8000000
    mov fp, #0
    bl _cstart
hang: b hang
```

// Disassembly of start.o

0000000 <_start>:

0: mov sp, #0x80000000
4: mov fp, #0
8: bl 0 <_cstart>

0000000c <hang>:

c: b c <hang>

Note: the address of _cstart is 0. Why?

It does know the address of hang. What's different between the two?

// Disassembly of clock.elf

<_start>:

8000:	mov	sp, 0x8000000
8004:	bl	80c8 <_cstart>

<hang>:

8008:	b	8008 <hang>
-------	---	-------------

Address of _cstart now known — 80c8
Who figured that out?

How is data handled?

Factors to consider:

extern vs static

read-only vs read-write

initialized vs uninitialized

```
// uninitialized  
int gNum;  
static int sgNum;  
  
// initialized  
int iNum = 1;  
static int siNum = 2;  
  
// const  
const int cNum = 3;  
static const int scNum = 4;
```

```
% arm-none-eabi-nm -S data.o
00000000 00000034 T binky
00000000 00000004 R cNum
00000004 00000004 C gNum
00000004 00000004 D iNum
00000034 0000001c T main
00000000 00000004 b sgNum
00000000 00000004 d siNum
```

The global uninitialized gNum is common (C).

The static const scNum was optimized out!

Guide to symbols (nm)

T/t - text (code)

D/d - read-write data

R/r - read-only data (**const**)

B/b - bss (*Block Started by Symbol*)

C - common (instead of B)

lowercase letter means static

SECTIONS

{

.text 0x8000 : { start.o(.text*)
 (.text) }

.data : { *(.data*) }

.rodata : { *(.rodata*) }

__bss_start__ = .;

.bss : { *(.bss*)
 *(COMMON) }

__bss_end__ = ALIGN(8);

}

Sections

Instructions go in `.text`

Data goes in `.data`

const data (read-only) goes in `.rodata`

Uninitialized data goes in `.bss`