

Memory Management

**Linking/Relocation
Loading
Starting
Heap**

Baremetal Section

Processor and memory architecture

Peripherals: GPIO, timers, UART

Assembly language and machine code

From C to assembly language

Functions and stack frames

Serial communication and strings

Modules and libraries: Linking

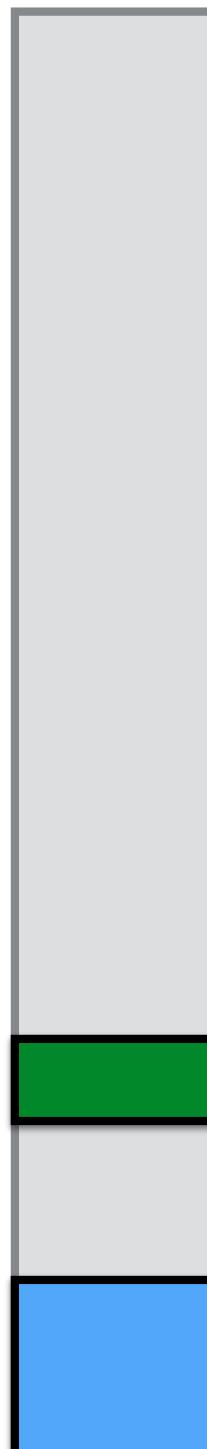
Memory Map: Linking, starting, and the heap

gpio
timer
uart
printf
malloc
keyboard
mailbox
fb
gl
console
monitor



The Memory Map

Memory Map



100000000_16

4 GB

Peripheral Registers

020000000_16

010000000_16

Memory Map



01000000_{16}

256 MB

Memory Map

GPU

10000000_{16}

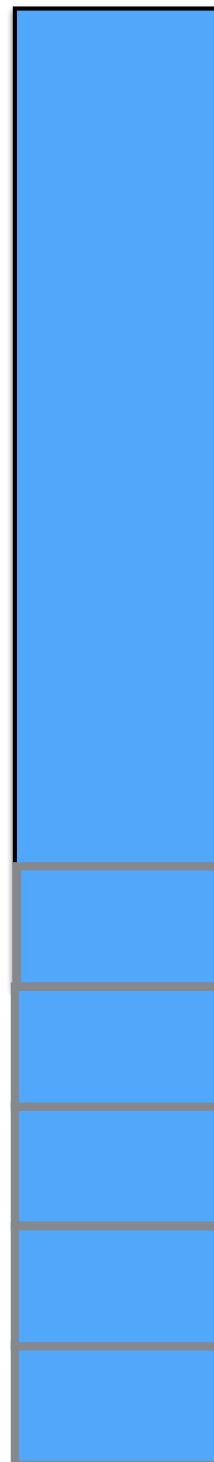
256 MB

CPU

08000000_{16}

128 MB

08000000₁₆



(uninitialized data) bss
(read-only data) rodata
data
text
interrupt vectors

00008000₁₆

linking/

```
// uninitialized global and static  
int i;  
static int j;
```

```
// initialized global and static  
int k = 1;  
static int l = 2;
```

```
// initialized global and static const  
const int m = 3;  
static const int n = 4;
```

```
// extern variable  
extern int p;
```

```
% make tricky.o  
% arm-none-eabi-nm -S -n tricky.o
```

```
U p  
00000000 00000004 b j  
00000000 00000004 D k  
00000000 00000004 R m  
00000000 00000044 T tricky  
00000004 00000004 C i  
00000004 00000004 d l
```

```
# The global uninitialized variable i  
# is in common (C).
```

```
# The static const variable n  
# has been optimized out.
```

Symbols

Types

- **extern (undefined)**
- **global vs static**
- **initialized vs uninitialized**
- **const vs non-const**

Guide to Symbols

T/t - text

D/d - (read-write) data

R/r - (read-only) data

B/b - bss (*Block Started by Symbol*)

C - common (used instead of B)

lower-case letter means static

Sections

Instructions go in .text

Data goes in .data

const data (read-only) goes in .rodata

Uninitialized data goes in .bss

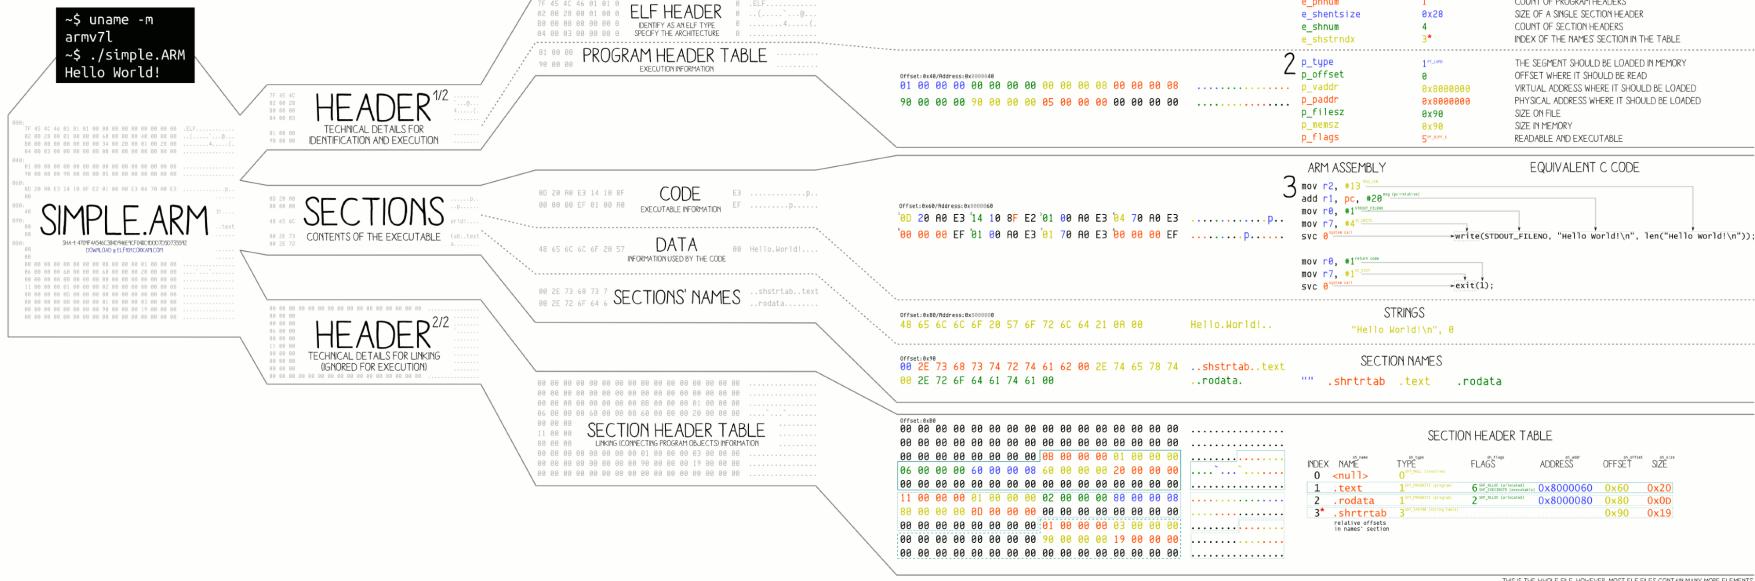
+ other information about the program

- symbols, relocation, debugging, ...**

ELF¹⁰¹ a Linux executable walkthrough

ANGE ALBERTINI
CORKAMIC.COM

DISSECTED FILE



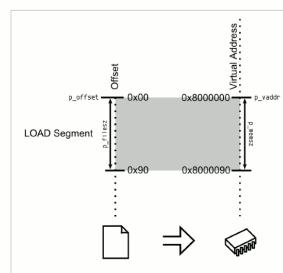
LOADING PROCESS

1 HEADER

THE ELF HEADER IS PARSED
THE PROGRAM HEADER IS PARSED
(SECTIONS ARE NOT USED)

2 MAPPING

THE FILE IS MAPPED IN MEMORY
ACCORDING TO ITS SEGMENT(S)



3 EXECUTION

ENTRY IS CALLED
SYSCALLS¹⁰¹ ARE ACCESSED VIA:
- SYSCALL NUMBER IN THE R7 REGISTER
- CALLING INSTRUCTION SVC

TRIVIA

THE ELF WAS FIRST SPECIFIED BY U.S. L. AND U.I.
FOR UNIX SYSTEM V, IN 1989

THE ELF IS USED, AMONG OTHERS, IN:

- LINUX, ANDROID, *BSD, SOLARIS, BEOS
- PSP, PLAYSTATION 2-4, DREAMCAST, GAMECUBE, WII
- VARIOUS OSES MADE BY SAMSUNG, ERICSSON, NOKIA,
- MICROCONTROLLERS FROM ATMEL, TEXAS INSTRUMENTS

memmap

```
.text : {
    start.o (.text)
    *(.text*)
    . = ALIGN(8);
} > ram
.data : { *(.data*) } > ram
.rodata : { *(.rodata*) } > ram
.bss : {
    __bss_start__ = .;
    *(.bss)
    *(COMMON)
    . = ALIGN(8);
    __bss_end__ = .;
} > ram
```

_start must be at #0x8000

```
% arm-none-eabi-nm -S main.exe | sort
00008000          T _start          # text
0000800c          t hang
00008010 00000038 T _cstart
00008048 00000048 T tricky
00008090 00000044 T main
000080d4 00000004 D k          # data
000080d8 00000004 d l
000080dc 00000004 D p
000080e0 00000004 R m          # rodata
000080e4          B __bss_start__ # bss
000080e4 00000004 b j
000080e8 00000004 B i
000080ec          B __bss_end__
```

```
% arm-none-eabi-size main.exe
text      data      bss      dec      hex filename
 212        12         8     232      e8 main.exe
```



```
% a-none-eabi-size *.o
text      data      bss      dec      hex filename
 56          0         0      56      38 cstart.o
 68          4         0      72      48 main.o
 12          0         0      12      c start.o
 76          8         4     88      58 tricky.o
```

size reports the size of the text

Note that the sum of the sizes of the .o's
is equal to the size of the main.exe

Relocation

```
// start.s
```

```
.globl _start
```

```
start:
```

```
    mov sp, #0x8000
```

```
    mov fp, #0
```

```
    bl _cstart
```

```
hang:
```

```
    b hang
```

// Disassembly of start.o (start.list)

0000000 <_start>:

0: mov sp, #0x8000
4: mov fp, #0
8: bl 0 <_cstart>

0000000c <hang>:

c: b c <hang>

// Note: the address of _cstart is 0

// Why?

// _start doesn't know where c_start is!

// Note it does know the address of hang

// Disassembly of main.exe.list

00008000 <_start>:

 8000: mov sp, #0x8000

 8004: mov fp, #0

 8008: bl 8010 <_cstart>

0000800c <hang>:

 800c: b 800c <hang>

00008010 <_cstart>:

 8010: push {r3, lr}

// Note: the address of _cstart is #8010

// Now _start knows where _cstart is!

Loading and Starting

Loading and Starting

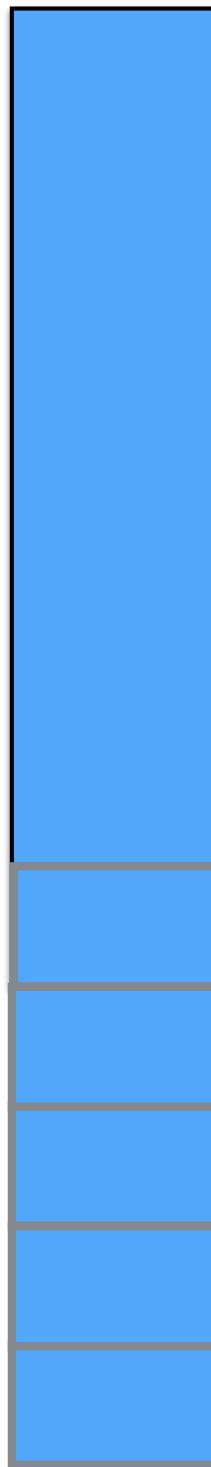
The boot loader copies the binary (*.bin) to memory starting at location #0x8000

Then the loader begins execution of the program by branching to #0x8000

Make sure that _start is at location #0x8000

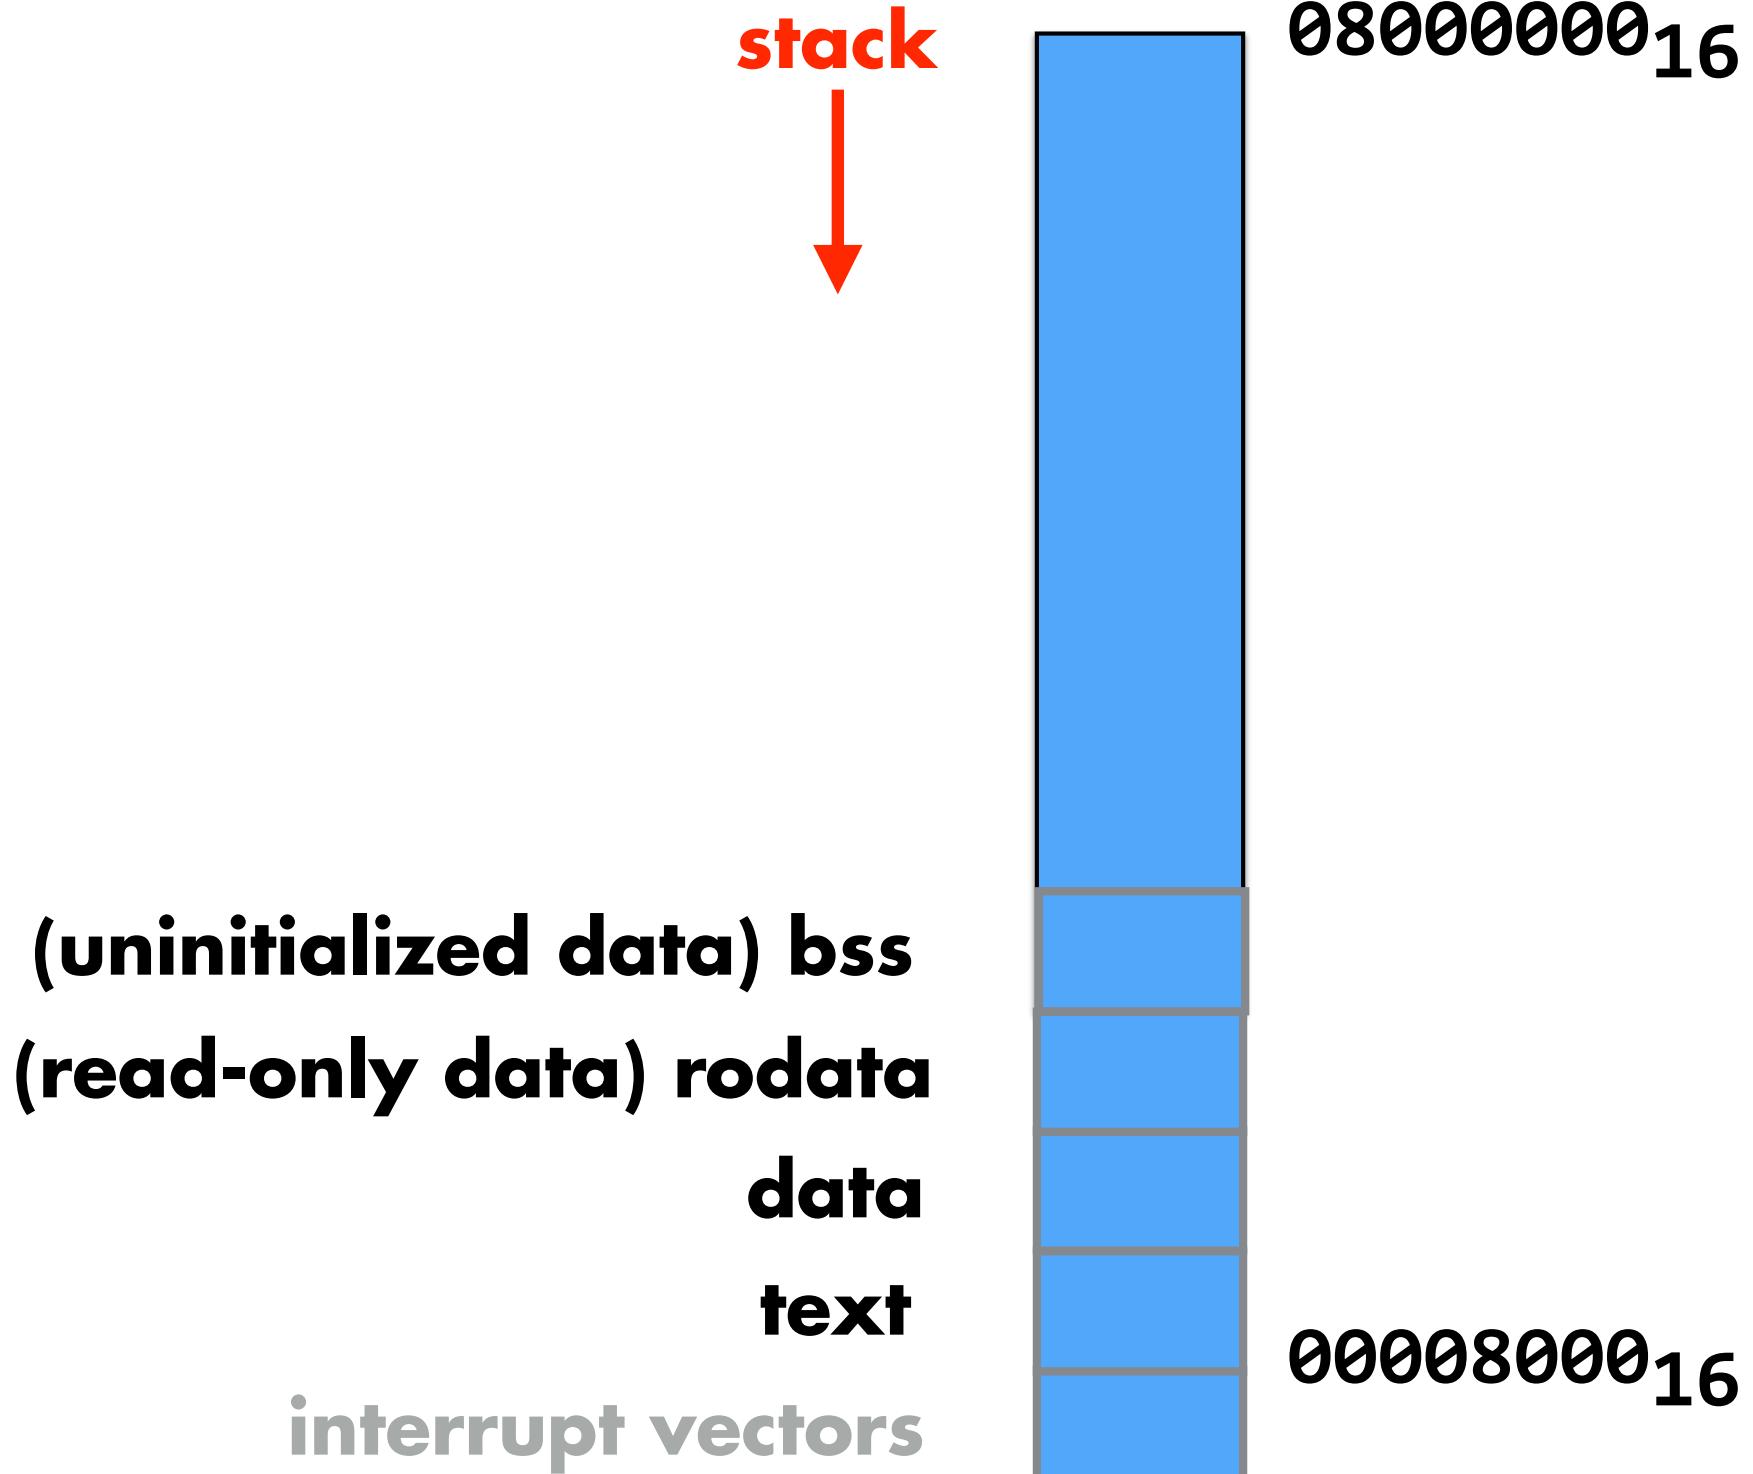
_start needs to setup the stack

08000000_{16}



00008000_{16}

(uninitialized data) bss
(read-only data) rodata
data
text
stack
↓



```
// better start.s
```

```
.globl _start
_start:
    mov sp, #0x08000000
    mov fp, #0
    bl _cstart
hang:
    b hang
```

```
// cstart.c - initializes bss to 0

extern int __bss_start__;
extern int __bss_end__;
void main();

void _cstart() {
    int* bss = &__bss_start__;
    int* bss_end = &__bss_end__;
    while( bss < bss_end )
        *bss++ = 0;
    main();
}
```

Memory Allocation and the Heap

Memory Allocation?

Compile-time memory allocation

Run-time memory allocation

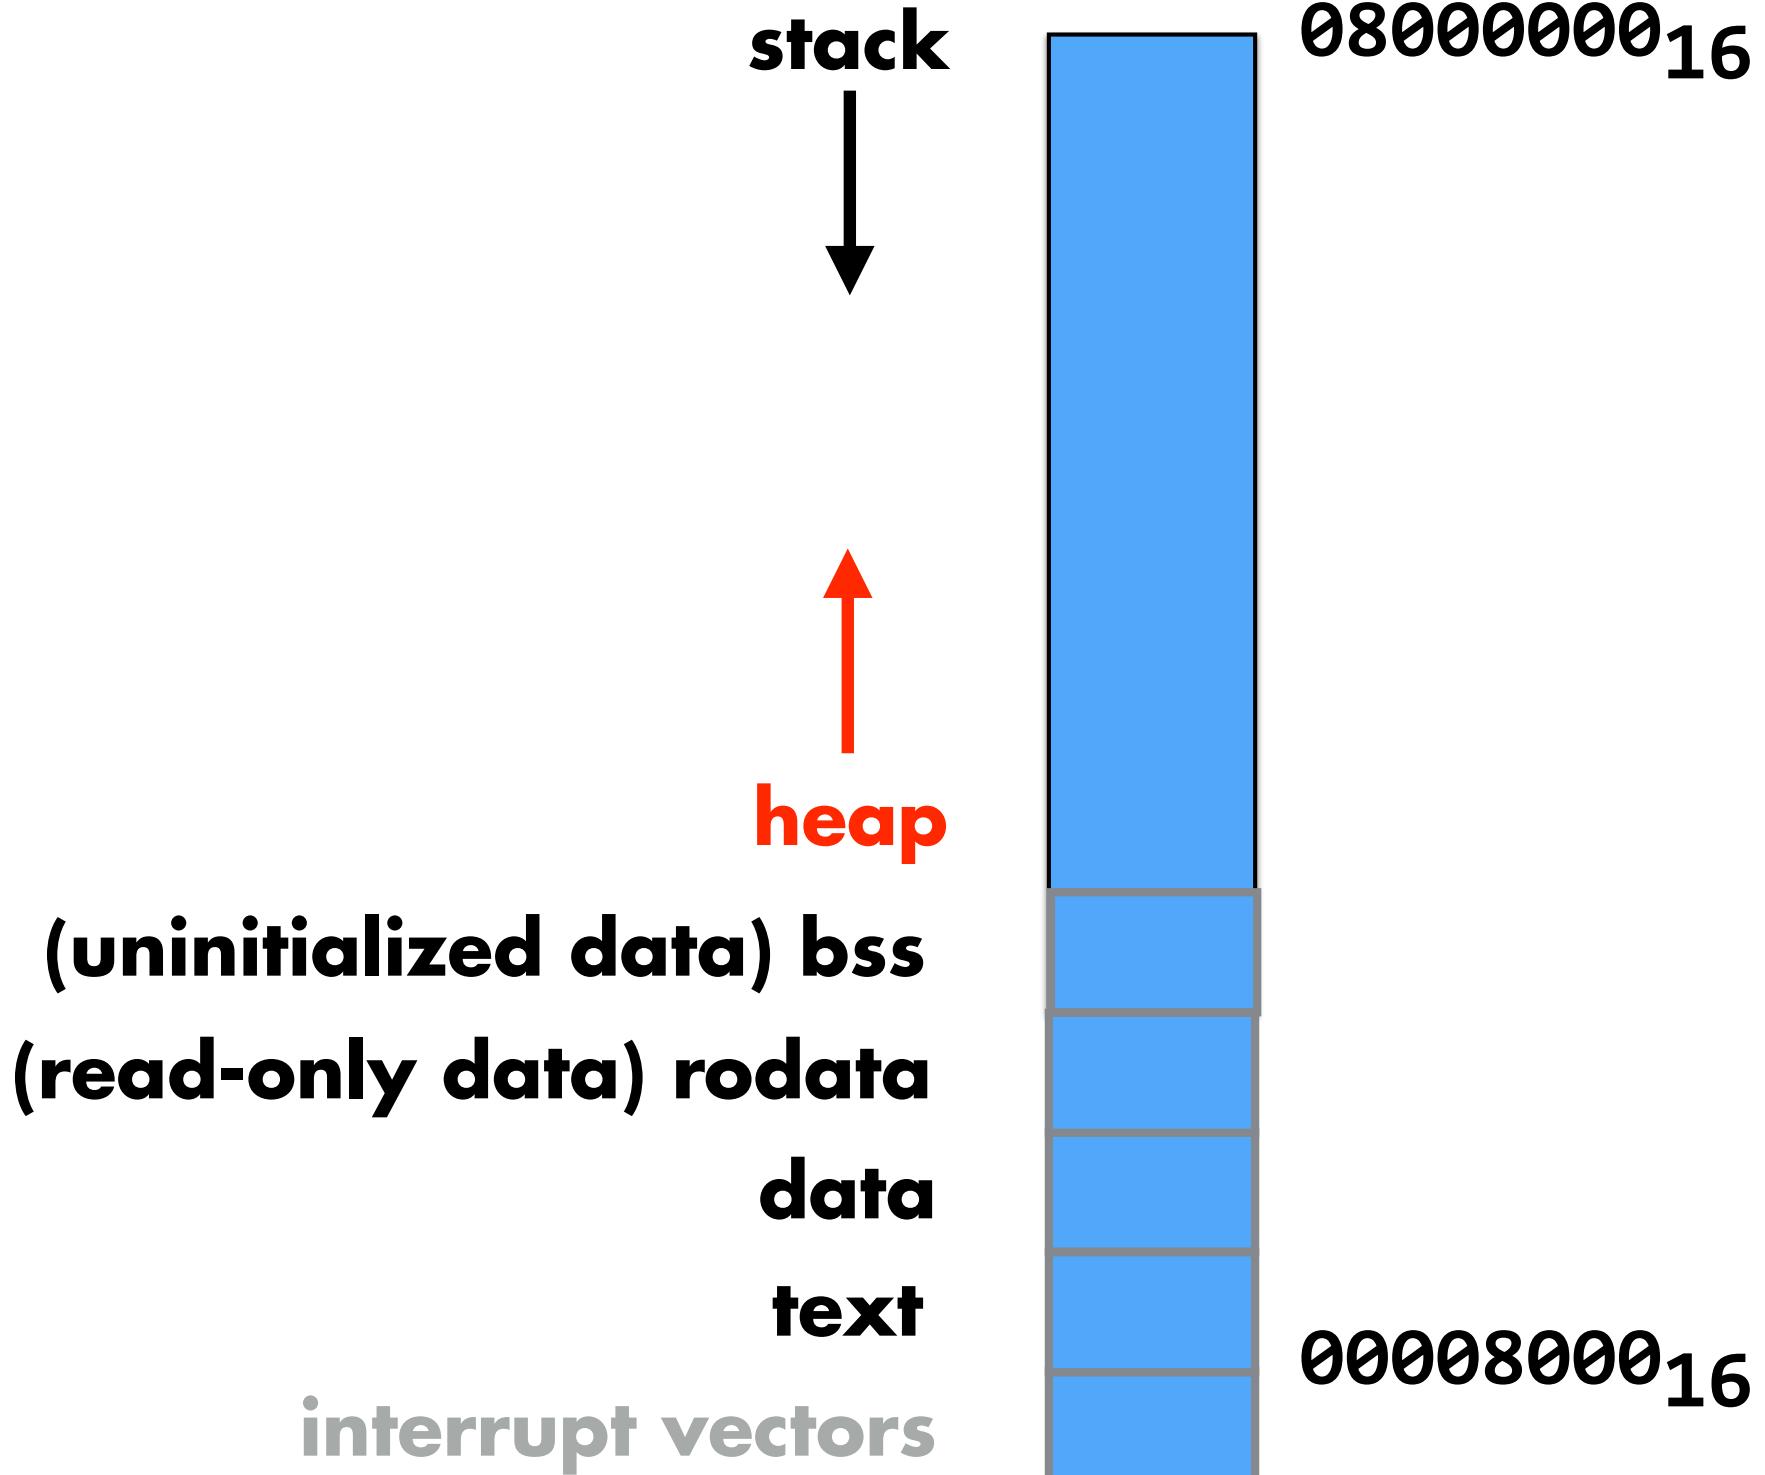
Why at run-time?

- 1. Don't know the size of an array when compiling**
- 2. Dynamic data structures such as lists and trees**

API

```
void *malloc( size_t size );  
void free( void *pointer );
```

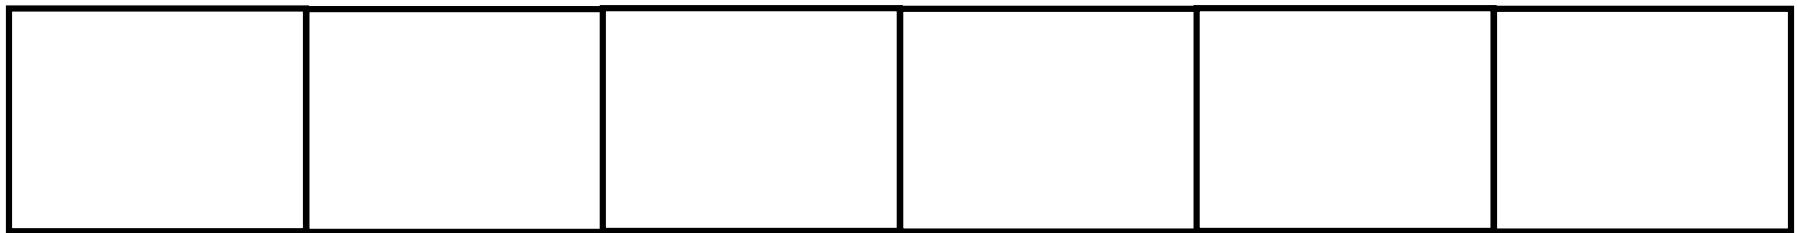
// Note that void* is a generic pointer



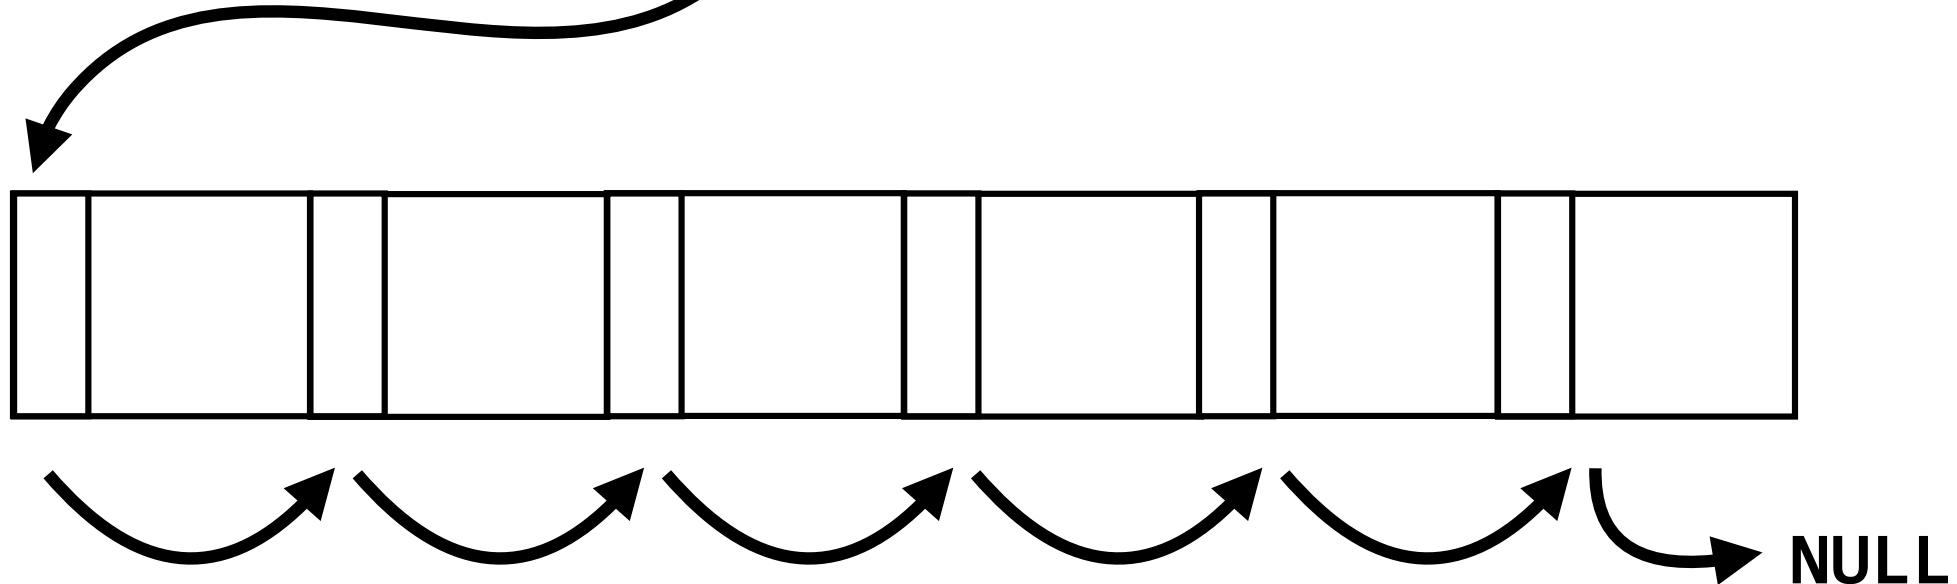
sbrk

list.c and block.c

```
newblock( nelements=6, nsize=16 );
```

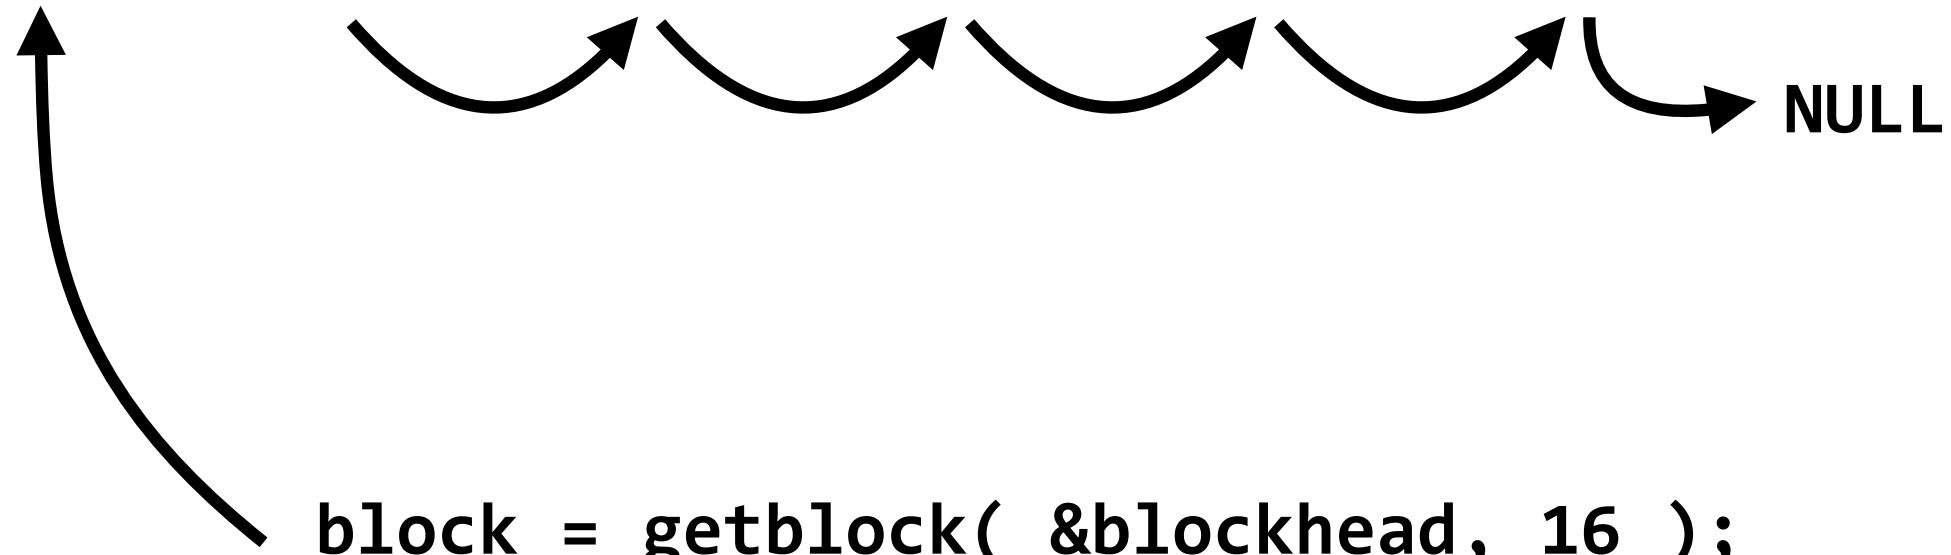
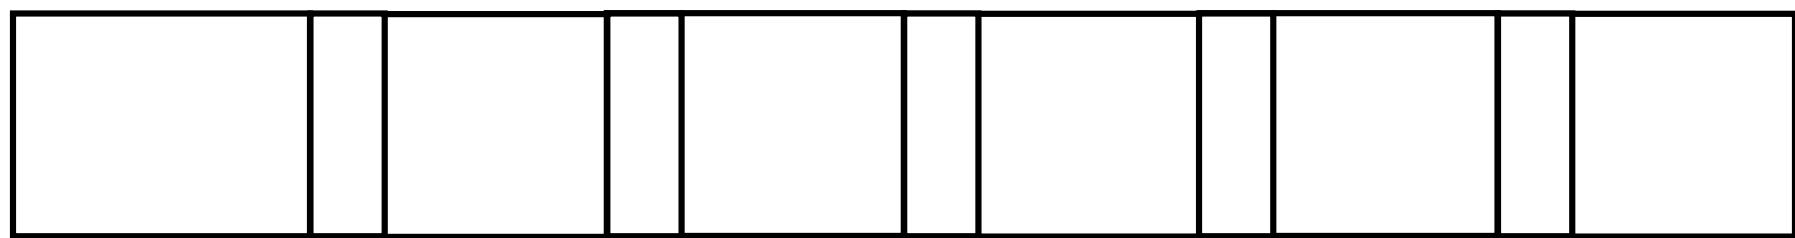


```
Block *blockhead = ;
```

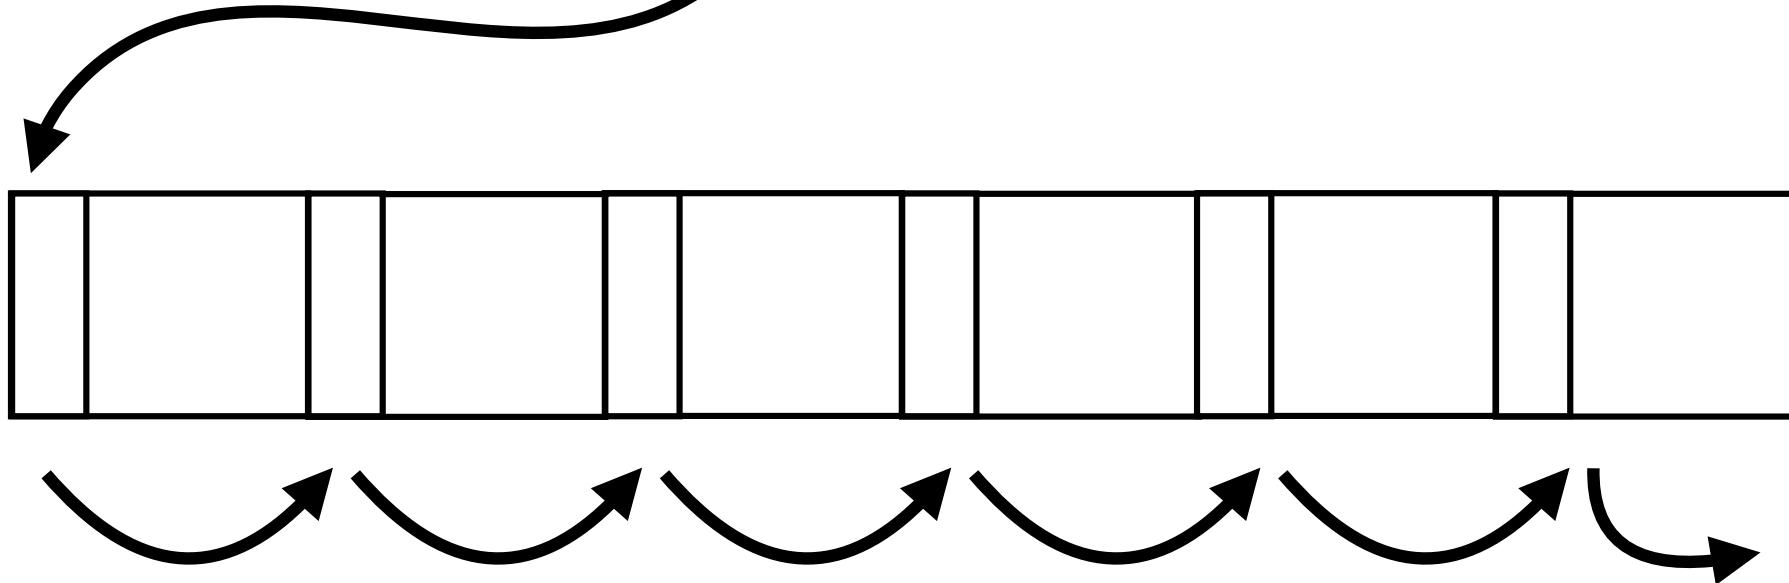


```
typedef struct b {  
    struct b *next;  
} Block;
```

```
Block *blockhead = ;
```



```
Block *blockhead = ;
```



```
getblock( &blockhead, block );
```

Variable Size malloc/free

just malloc is easy

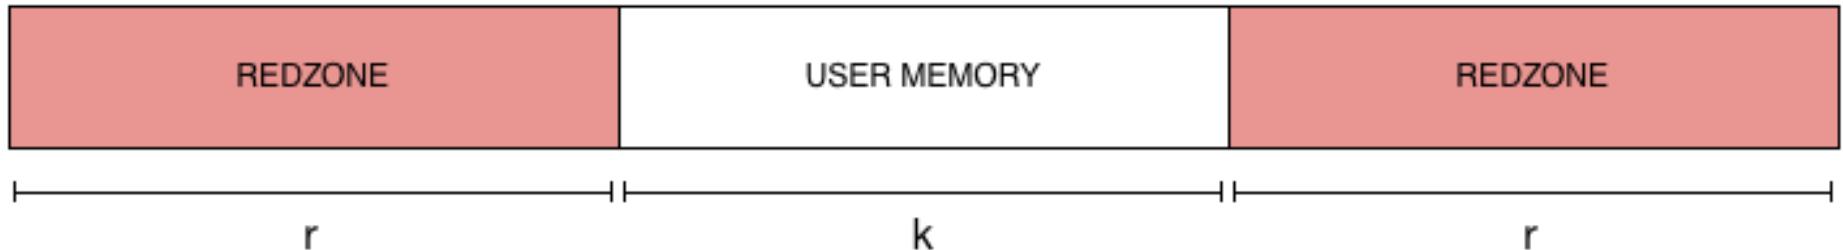


malloc with free is hard



- **free returns blocks that can be re-allocated**
- **malloc should search to see if there is a block of sufficient size. Which block should it choose (best-fit, first-fit, largest)?**
- **malloc may use only some of the block. It splits the block into two sub-blocks of smaller sizes**
- **splitting blocks causes fragmentation**

Memory Corruption



**Write special value (#0xDEADBEEF) to red zone
Look for unintended writes to the red zones**

red zone malloc
Assignment 4