

Goals for today

Thanks for the memory!

Linker memory map

Address space layout

Loading


How an executable file becomes a running program

Heap allocation

Malloc, realloc, and free

Admin

printf perseverance and pride!!



How is global data handled?

Factors to consider:

extern vs static

read-only vs read-write

initialized vs uninitialized

// uninitialized

```
int gNum;  
static int sgNum;
```

// initialized

```
int iNum = 1;  
static int siNum = 2;
```

// const

```
const int cNum = 3;  
static const int scNum = 4;
```

Note: In C, uninitialized global variables are zeroed!

```
% arm-none-eabi-nm -S data.o
00000000 00000034 T binky
00000000 00000004 R cNum
00000004 00000004 C gNum
00000004 00000004 D iNum
00000034 0000001c T main
00000000 00000004 b sgNum
00000000 00000004 d siNum
```

The global uninitialized gNum is common (C).

The static const scNum was optimized away!

Guide to symbols (nm)

T/t - text (code)

D/d - read-write data

R/r - read-only data (const)

B/b - bss (*Block Started by Symbol*)

C - common (instead of B)

lowercase for static, uppercase for extern

memmap linker script

SECTIONS

```
{  
    .text 0x8000 : { start.o(.text*)  
                    *(.text*) }  
    .data :        { *(.data*) }  
    .rodata :      { *(.rodata*) }  
  
    __bss_start__ = .;  
    .bss :          { *(.bss*)  
                    *(COMMON) }  
    __bss_end__ = ALIGN(8);  
}
```

```
// memmap wraps bfs section with these symbols
extern int __bss_start__, __bss_end__;


extern void main(void);

void _cstart(void) {
    int *bss = &__bss_start__;
    int *bss_end = &__bss_end__;

    while (bss < bss_end) {
        *bss++ = 0;
    }

    main();
}
```

Memory Map



SECTIONS

```
{
    .text 0x8000 : { start.o(.text*)
                  { *(.text*) }
    .data :      { *(.data*) }
    .rodata :    { *(.rodata*) }

    __bss_start__ = .;
    .bss :        { *(.bss*)
                  { *(COMMON) }
    __bss_end__ = ALIGN(8);
}
```

(zeroed data) bss

(read-only data) rodata

data

text

interrupt vectors



0x80000000

_start:

mov sp, #0x80000000

mov fp, #0

bl _cstart

__bss_end__

blink.bin
0x8000

Global allocation

- + **Convenient, albeit hacky**

 - Global scope, all can access

 - No encapsulation, hard to track use/dependencies

- **Size fixed at declaration, no option to resize**

- +/- **Scope/lifetime is global/whole program**

 - One shared namespace, manually avoid conflicts

Stack allocation

- + **Efficient**

 - Fast to allocate/deallocate, ok to oversize

- + **Convenient**

 - Automatic alloc/dealloc on function entry/exit

 - Can assign with static initializer

- + **Reasonable type safety**

- **Size fixed at declaration, no option to resize**

 - Cannot reassign array — there is no pointer!

- +/- **Scope/lifetime dictated by control flow**

Heap allocation

- + **Moderately efficient**

 - Have to search for available space, update record-keeping

- + **Very plentiful**

 - Heap enlarges on demand to limits of address space

- + **Versatile, under programmer control**

 - Can precisely determine scope, lifetime

 - Can be resized

- **Lots of opportunity for error**

 - Low type safety

- **Heap memory errors**

 - (allocate wrong size, use after free, double free)

- **Leaks** (much less critical)

Heap interface

```
void *malloc(size_t nbytes);  
void free(void *ptr);  
void *realloc(void *ptr, size_t nbytes);
```

void* pointer

Variable of type address with unspecified/unknown pointee type

What you can do with a void *

Pass to/from function, pointer assignment

What you cannot

Cannot dereference

Cannot do pointer arithmetic

Cannot use array indexing (both arithmetic & dereference!)

Why do we need a help?

Let's see an example!

`code/heap.c`

How is a heap implemented?



Tracing the heap

```
p1 = malloc(4)
```



```
p2 = malloc(5)
```




```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```




How are we going to do that??

Bump Memory Allocator


malloc.c

Now: track inuse, recycle memory

Standard technique is to track status (length and free/inuse) about each block in the word preceding the payload data.



The heap now has an implicit list to walk all the blocks



Heap implementation design

just `malloc` is easy 😊
`malloc` with `free` is hard 😞

`free` recycles blocks to service later requests

`malloc` search for existing unused block. Which block to choose (best-fit, first-fit)?

`malloc` may need only some of the block and may split it. Splitting blocks causes fragmentation.

`free` can coalesce with neighboring free blocks to reduce fragmentation

Points to ponder

What happens if you forget to free a pointer after you are done using it?

Can you refer to a pointer after it has been freed?

What is stored in the memory that you malloc?

What if you free a pointer that you didn't malloc?

Can you free the same pointer twice?

Wouldn't it be nice to not have to worry about freeing memory?